

Evaluation Report

****Document ID:**** evaluation-v1

****Date:**** 2026-08-31

1. Executive Summary

This report presents the evaluation of the recommendation model, LLM judge, and uplift targeting com

Key Findings

2. Recommendation Model Evaluation

2.1 Operating Characteristics

The model produces `score\_live` at the operating threshold of 0.62.

2.2 Calibration

2.3 Decision Curve Analysis

Net benefit analysis shows the model provides benefit across a range of threshold probabilities, but

- Unknown true outcome in synthetic data
- No independent validation set

2.4 Leakage Assessment

3. LLM Judge Evaluation

3.1 Comparison with Clinical Labels

The LLM judge scores were compared against the blinded clinical labels from the two reviewers.

3.2 Agreement Metrics

3.3 Limitations

****Critical:**** The LLM judge shares a model family with the recommendation system. This prevents use

****Status:**** NOT APPROVED FOR AUTONOMOUS SCORING

The judge may be described only as triage support. Unless lower adjusted bounds for sensitivity and

4. Uplift Targeting Evaluation

4.1 Reported Metrics

4.2 Evaluation Limitations

The `model\_auc\_reported` field cannot validate causal uplift because:

1. No causal assignment contract exists
2. Uplift requires counterfactual outcomes
3. Selection into targeting is not random

4.3 Diagnostics Performed

- Cross-fitted Qini/AUUC (not estimable without causal assignment)
- Uplift calibration (not estimable)
- Doubly robust policy value (not estimable)

****Status:**** CANNOT VALIDATE CAUSAL UPLIFT

5. Label Quality Assessment

5.1 Reviewer Agreement

Two reviewers labeled 120 pairs. Their agreement was assessed before looking at judge scores.

5.2 Adjudication Status

****Status:**** NO ADJUDICATION EXISTS

The two reviewers were not adjudicated. Without third-clinician adjudication of disagreements, no si

6. Model Subgroup Performance

6.1 Site Stratification

6.2 Service Stratification

Subgroup estimates are suppressed when fewer than 5 service clusters or 20 episodes exist.

6.3 Demographic Stratification

7. Sensitivity Analyses

7.1 Time Horizon

7.2 Data Source

7.3 Alternate Specifications

8. Authorization and Access

8.1 Model Access

8.2 Authorization Violations

Zero authorization violations were detected in testing. However, this does not prove against the uns

9. Patient-Level Reliability

9.1 Intra-Patient Correlation

Patients with multiple episodes show intraclass correlation. This affects standard error estimation.

9.2 Clustering by Clinician

Clinician-level clustering is accounted for in the inference procedure.

10. Conclusions and Recommendations

10.1 Recommendation Model

The model shows reasonable calibration but lacks independent validation. Use for decision support on

10.2 LLM Judge

****NOT APPROVED**** for autonomous scoring due to same-family dependence. May be used as triage support

10.3 Uplift Targeting

Cannot validate causal uplift without causal assignment contract. Current metrics are descriptive on

10.4 Labels

No ground truth exists due to lack of adjudication. Current labels are expert opinion only.

Appendix A: Tables

Table A1: Model Performance by Service

[See machine tables]

Table A2: Judge Agreement by Service

[See machine tables]

Table A3: Subgroup Sensitivity

[See machine tables]

Appendix B: Figures

Figure B1: Calibration Curve

[Attached]

Figure B2: Decision Curve

[Attached]

Figure B3: Score Distribution

[Attached]

This evaluation is based on synthetic data. All limitations are explicitly stated.